

## HW6


### Question 1

In this problem you, can simulate a simplified airport security system at a busy airport. Passengers arrive according to a Poisson distribution with  $\lambda_1 = 5$  per minute (i.e., mean interarrival rate  $\mu_1 = 0.2$  minutes) to the ID/boarding-pass check queue, where there are several servers who each have exponential service time with mean rate  $\mu_2 = 0.75$  minutes. [Hint: model them as one block that has more than one resource.] After that, the passengers are assigned to the shortest of the several personal-check queues, where they go through the personal scanner (time is uniformly distributed between 0.5 minutes and 1 minute).

Use the Arena software (PC users) or Python with SimPy (Mac users) to build a simulation of the system, and then vary the number of ID/boarding-pass checkers and personal-check queues to determine how many are needed to keep average wait times below 15 minutes.

Answer:

First, try 3 servers at ID & boarding pass check, then 2 personal scanner lines with one server each. Run 10 replications with 5 hours each. The result for avg waiting time is 0.44hr (26.4 min). And the personal scanner line each need to wait about 0.31hr (18.6min), which is already higher than 15 min.



Create

Name:  Entity Type:

Time Between Arrivals

Type:  Expression:  Units:

Entities per Arrival:  Max Arrivals:  First Creation:

OK Cancel Help

Process

Name:  Type:

Logic

Action:  Priority:

Resources:

Add... Edit... Delete

Delay Type:  Units:  Allocation:

Expression:

☒ Report Statistics

OK Cancel Help

Decide

Name:  Type:

If:

Value:

OK Cancel Help

Process

Name:  Type:

Logic

Action:  Priority:

Resources:

Add... Edit... Delete

Delay Type:  Units:  Allocation:

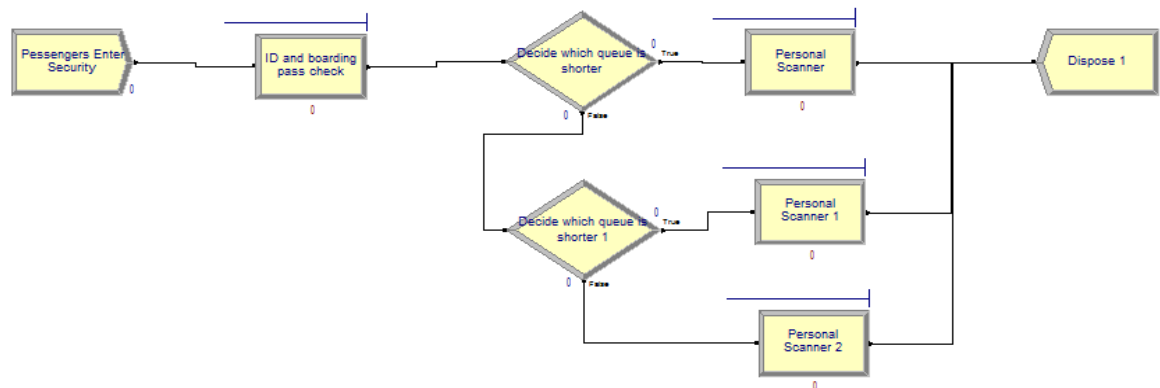
Expression:

☒ Report Statistics

OK Cancel Help

		Waiting Time	
		Average	
Wait Time	Average	ID and boarding pass check.Queue	0.1300
Entity 1	0.4396	Personal Scanner 1.Queue	0.3108
		Personal Scanner.Queue	0.3114

Second, try to add one personal scanner. The result of avg waiting time is 0.31hr(18.6min). It seems that the waiting time at ID and boarding pass check is too long.



		Waiting Time	
		Average	
Wait Time	Average	ID and boarding pass check.Queue	0.2636
Entity 1	0.3156	Personal Scanner 1.Queue	0.05120456
		Personal Scanner 2.Queue	0.04611909
		Personal Scanner.Queue	0.05924787

Third, try to add additional one server at ID & boarding pass check. The result of avg waiting time is 0.275hr(16.5 min)

		Waiting Time	
		Average	
Wait Time	Average	ID and boarding pass check.Queue	0.04396864
Entity 1	0.2756	Personal Scanner 1.Queue	0.2308
		Personal Scanner 2.Queue	0.2287
		Personal Scanner.Queue	0.2385

Forth, try to add additional one personal scanner. The result of avg waiting time is 0.1088hr (6.528min), which is less than 15 min. Great! Then we use 4 servers at ID check, and use 4 personal scanners.

		Waiting Time	
		Average	
Wait Time	Average	ID and boarding pass check.Queue	0.07457095
Entity 1	0.1088	Personal Scanner 1.Queue	0.03570406
		Personal Scanner 2.Queue	0.02793915
		Personal Scanner 3.Queue	0.02387695
		Personal Scanner.Queue	0.04553027

## Question 2

The breast cancer data set at <http://archive.ics.uci.edu/ml/machine-learning-databases/breast-cancer-wisconsin/> (description at <http://archive.ics.uci.edu/ml/datasets/Breast+Cancer+Wisconsin+%28Original%29>) has missing values.

1. Use the mean/mode imputation method to impute values for the missing data.
2. Use regression to impute values for the missing data.
3. Use regression with perturbation to impute values for the missing data.
4. (Optional) Compare the results and quality of classification models (e.g., SVM, KNN) build using (1) the data sets from questions 1,2,3; (2) the data that remains after data points with missing values are removed; and (3) the data set when a binary variable is introduced to indicate missing values.

Answer:

2.1

# First, read the data

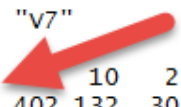
```
data <- read.table("breast-cancer-wisconsin.data.txt", stringsAsFactors = FALSE, header = FALSE, sep=",")
```

# Try to find the missing value

```
for (i in 2:11){  
  print(paste0("V",i))  
  print(table(data[,i]))  
}
```

```
[1] "V7"
```

```
? 10 2 3 4 5 6 7 8 9  
16 402 132 30 28 19 30 4 8 21 9
```



⇒ Only V7 with missing 16 values (show in question marks)

# show the missing data

```
data[which(data$V7=="?"),]
```

	v1	v2	v3	v4	v5	v6	v7	v8	v9	v10	v11
24	1057013	8	4	5	1	2	?	7	3	1	4
41	1096800	6	6	6	9	6	?	7	8	1	2
140	1183246	1	1	1	1	1	?	2	1	1	2
146	1184840	1	1	3	1	2	?	2	1	1	2
159	1193683	1	1	2	1	3	?	1	1	1	2
165	1197510	5	1	1	1	2	?	3	1	1	2
236	1241232	3	1	4	1	2	?	3	1	1	2
250	169356	3	1	1	1	2	?	3	1	1	2
276	432809	3	1	3	1	2	?	2	1	1	2
293	563649	8	8	8	1	2	?	6	10	1	4
295	606140	1	1	1	1	2	?	2	1	1	2
298	61634	5	4	3	1	2	?	2	3	1	2
316	704168	4	6	5	6	7	?	4	9	1	2
322	733639	3	1	1	1	2	?	3	1	1	2
412	1238464	1	1	1	1	1	?	2	1	1	2
618	1057067	1	1	1	1	1	?	1	1	1	2


# calculate the missing %

```
nrow(data[which(data$V7=="?"),])/nrow(data)
```

```
[1] 0.02288984
```

⇒ Smaller than 5%, okay to use imputation. In this case, v7 is not continuous number, so should use mode instead of mean to impute. The mode for V7 should be 1. Then use 1 to impute missing values.

```
[1] "v7"
```



	?	1	10	2	3	4	5	6	7	8	9
	16	402	132	30	28	19	30	4	8	21	9

## 2.2

# Not to include the response variable in regression imputation

```
data_modified <- data[-missing, 2:10]
```

```
data_modified$V7 <- as.integer(data_modified$V7)
```

# run linear regression

```
model <- lm(V7~V2+V3+V4+V5+V6+V7+V8+V9+V10, data_modified)
```

```
summary(model)
```

```
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -0.616652   0.194975  -3.163  0.00163 **
V2           0.230156   0.041691   5.521 4.83e-08 ***
V3          -0.067980   0.076170  -0.892  0.37246
V4           0.340442   0.073420   4.637 4.25e-06 ***
V5           0.339705   0.045919   7.398 4.13e-13 ***
V6           0.090392   0.062541   1.445  0.14883
V8           0.320577   0.059047   5.429 7.91e-08 ***
V9           0.007293   0.044486   0.164  0.86983
V10          -0.075230   0.059331  -1.268  0.20524
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 2.274 on 674 degrees of freedom
Multiple R-squared:  0.615,    Adjusted R-squared:  0.6104
F-statistic: 134.6 on 8 and 674 DF,  p-value: < 2.2e-16
```

```
step(model)
```

```
call:
lm.default(formula = V7 ~ V2 + V4 + V5 + V8, data = data_modified)
```

```
Coefficients:
(Intercept)          V2          V4          V5          V8
   -0.5360      0.2262      0.3173      0.3323      0.3238
```

⇒ V2, V4, V5, V8 are important variables to predict V7

```
model2 <- lm(V7~V2+V4+V5+V8, data_modified)
```

```
summary(model2)
```

```
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept) -0.53601    0.17514  -3.060  0.0023 **
V2           0.22617    0.04121   5.488 5.75e-08 ***
V4           0.31729    0.05086   6.239 7.76e-10 ***
V5           0.33227    0.04431   7.499 2.03e-13 ***
V8           0.32378    0.05606   5.775 1.17e-08 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' '

Residual standard error: 2.274 on 678 degrees of freedom
Multiple R-squared:  0.6129,    Adjusted R-squared:  0.6107
F-statistic: 268.4 on 4 and 678 DF,  p-value: < 2.2e-16
```

⇒ All variables are significant, use this model to predict V7

# predict(impute) V7

```
V7_hat <- predict(model2, newdata = data[missing,])
```

```
> v7_hat
      24      41      140      146      159      165      236      250      276
5.4585352 7.9816106 0.9872832 1.6218560 0.9807851 2.2157441 2.7152652 1.7634059 2.0741942
      293      295      298      316      322      412      618
6.0866099 0.9872832 2.5265324 5.2438347 1.7634059 0.9872832 0.6634986
```

```
data_reg_imp <- data
data_reg_imp[missing,]$V7 <- V7_hat
data_reg_imp$V7 <- as.numeric(data_reg_imp$V7)
```

2.3

```
V7_hat_pert <- rnorm(length(missing), V7_hat, sd(V7_hat))
```

```
V7_hat_pert
```

```
· v7_hat_pert
[1] 4.0777220 8.3863924 -0.8545876 5.1381323 1.7070775 0.4072891 3.7896436 3.3908019
[9] 3.3433164 5.4134808 4.3195118 3.3858147 3.8745124 -3.1181778 3.4668265 0.5644572
```

```
data_reg_pert_imp <- data
data_reg_pert_imp[missing,]$V7 <- V7_hat_pert
data_reg_pert_imp$V7 <- as.numeric(data_reg_pert_imp$V7)
```

### Question 3

Describe a situation or problem from your job, everyday life, current events, etc., for which optimization would be appropriate. What data would you need?

Answer:

Use optimization to allocate my time properly to maximize my happiness. (Happiness here is defined in a broad view. For example “rest” is a kind of happiness that can be get by doing activities such as watch TV, sleep or meditation) Following is an optimization model which specify the required data.

n activities

m happiness

$a_{ij}$  = amount of happiness  $j$  per unit of activity  $i$

$m_j$  = minimum daily required of happiness  $j$

$M_j$  = Maximum daily required of happiness  $j$

Variables:

$X_i$  = How many hour spend on activity  $i$  daily

Constraints:

$\sum_i a_{ij} X_i \geq m_j$  for each kind of happiness  $j$  (For example, it is must to have nutrition and have rest every day. Therefore set a minimum requirement)

$\sum_i a_{ij} X_i \leq M_j$  for each kind of happiness  $j$  (For example, rest too much is also not good. Therefore set

a Max value )

$\sum_i X_i \leq 24$  hours (A day only have 24 hours)

$X_i \geq 0$

Objective function:

Maximize  $\sum_i X_i \cdot a_{ij}$